


✓ Importing Twitter Data

```
import pandas as pd
import os
from google.colab import drive
drive.mount('/content/drive')
extracted_dir_path = '/content/drive/MyDrive/TwitterData'
# Load train and test data
train_data_path = os.path.join(extracted_dir_path, 'train.csv')
test_data_path = os.path.join(extracted_dir_path, 'test.csv')
# Load train and test data with specified encoding
train_data = pd.read_csv(train_data_path, encoding='ISO-8859-1')
test_data = pd.read_csv(test_data_path, encoding='ISO-8859-1')
```

 Mounted at /content/drive

✓ Data Preprocessing

```
# Select necessary columns in train data
train_data = train_data[['text', 'sentiment']]

# Select necessary columns in test data
test_data = test_data[['text', 'sentiment']]

# Remove null values if any
train_data = train_data.dropna()
test_data = test_data.dropna()
```

```
# Lowercase the text
train_data['text'] = train_data['text'].str.lower()
test_data['text'] = test_data['text'].str.lower()

# Remove punctuation and special characters
train_data['text'] = train_data['text'].str.replace('[^\w\s]', '')
test_data['text'] = test_data['text'].str.replace('[^\w\s]', '')

# Remove stopwords
import nltk
from nltk.corpus import stopwords

nltk.download('stopwords')
stop_words = set(stopwords.words('english'))

train_data['text'] = train_data['text'].apply(lambda x: ' '.join([word for word in x.split() if word not in stop_words]))
test_data['text'] = test_data['text'].apply(lambda x: ' '.join([word for word in x.split() if word not in stop_words]))

# Tokenize the text
from tensorflow.keras.preprocessing.text import Tokenizer

tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_data['text'])

X_train = tokenizer.texts_to_sequences(train_data['text'])
X_test = tokenizer.texts_to_sequences(test_data['text'])

# Pad sequences
from tensorflow.keras.preprocessing.sequence import pad_sequences

max_len = max([len(x) for x in X_train + X_test])
X_train = pad_sequences(X_train, maxlen=max_len, padding='post')
X_test = pad_sequences(X_test, maxlen=max_len, padding='post')
```

```
➞ [nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
```

```

from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
train_data['sentiment'] = label_encoder.fit_transform(train_data['sentiment'])

# Convert sentiment labels to numeric labels for test data
test_data['sentiment'] = label_encoder.transform(test_data['sentiment'])

# Separate features (X) and target (y) for train and test data
X_train = tokenizer.texts_to_sequences(train_data['text'])
X_train = pad_sequences(X_train, maxlen=max_len, padding='post')
y_train = train_data['sentiment']

X_test = tokenizer.texts_to_sequences(test_data['text'])
X_test = pad_sequences(X_test, maxlen=max_len, padding='post')
y_test = test_data['sentiment']

```

✓ LSTM Model

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense

# Define the model
lstm_model = Sequential([
    Embedding(input_dim=len(tokenizer.word_index) + 1, output_dim=100, input_length=max_len),
    LSTM(128),
    Dense(3, activation='softmax')
])

# Compile the model
lstm_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

print(" ")
print("LSTM Model Architecture : ")
print(lstm_model.summary())
print(" ")

# Train the model
lstm_history = lstm_model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))

```



LSTM Model Architecture :
Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 31, 100)	2659400
lstm (LSTM)	(None, 128)	117248
dense (Dense)	(None, 3)	387

=====
Total params: 2777035 (10.59 MB)
Trainable params: 2777035 (10.59 MB)
Non-trainable params: 0 (0.00 Byte)

None

Epoch 1/20
687/687 [=====] - 36s 49ms/step - loss: 0.9443 - accu
Epoch 2/20
687/687 [=====] - 6s 9ms/step - loss: 0.6443 - accur
Epoch 3/20
687/687 [=====] - 5s 7ms/step - loss: 0.4456 - accur
Epoch 4/20
687/687 [=====] - 4s 5ms/step - loss: 0.3134 - accur
Epoch 5/20
687/687 [=====] - 5s 7ms/step - loss: 0.2313 - accur
Epoch 6/20
687/687 [=====] - 4s 6ms/step - loss: 0.1754 - accur
Epoch 7/20
687/687 [=====] - 4s 6ms/step - loss: 0.1407 - accur
Epoch 8/20
687/687 [=====] - 5s 7ms/step - loss: 0.1208 - accur
Epoch 9/20
687/687 [=====] - 5s 7ms/step - loss: 0.1060 - accur
Epoch 10/20
687/687 [=====] - 7s 10ms/step - loss: 0.0866 - accu
Epoch 11/20
687/687 [=====] - 8s 11ms/step - loss: 0.0748 - accu
Epoch 12/20
687/687 [=====] - 7s 10ms/step - loss: 0.0705 - accu
Epoch 13/20
687/687 [=====] - 7s 11ms/step - loss: 0.0599 - accu
Epoch 14/20
687/687 [=====] - 7s 10ms/step - loss: 0.0521 - accu
Epoch 15/20
687/687 [=====] - 6s 8ms/step - loss: 0.0486 - accur
Epoch 16/20

```

687/687 [=====] - 4s 6ms/step - loss: 0.0441 - accuracy: 0.6505
Epoch 17/20
687/687 [=====] - 7s 10ms/step - loss: 0.0414 - accuracy: 0.6505
Epoch 18/20
687/687 [=====] - 5s 7ms/step - loss: 0.0386 - accuracy: 0.6505
Epoch 19/20
687/687 [=====] - 6s 9ms/step - loss: 0.0331 - accuracy: 0.6505
Epoch 20/20
687/687 [=====] - 5s 8ms/step - loss: 0.0338 - accuracy: 0.6505

```

```

# Evaluate the model
lstm_loss, lstm_accuracy = lstm_model.evaluate(X_test, y_test)
print("LSTM Test Accuracy:", lstm_accuracy)

```

```

➞ 111/111 [=====] - 1s 5ms/step - loss: 2.1022 - accuracy: 0.6505
LSTM Test Accuracy: 0.6505376100540161

```

✓ RNN Model

```

from keras.models import Sequential
from keras.layers import Embedding, SimpleRNN, Dense, Activation, Flatten
from keras.preprocessing.sequence import pad_sequences

# Pad sequences
X_train = pad_sequences(X_train, maxlen=max_len)
X_test = pad_sequences(X_test, maxlen=max_len)

# Build the RNN model
rnn_model = Sequential()

# Adjust the embedding input dimension based on your vocabulary size
vocab_size = len(tokenizer.word_index) + 1

rnn_model.add(Embedding(vocab_size, 32, input_length=max_len))
rnn_model.add(SimpleRNN(16, return_sequences=False, activation="relu"))
rnn_model.add(Dense(3)) # Three classes: negative, neutral, positive
rnn_model.add(Activation("softmax")) # Softmax for multi-class classification

# Compile the model
rnn_model.compile(loss="sparse_categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

print(" ")

```

```

print("RNN Model Architecture : ")
print(rnn_model.summary())
print(" ")

# Train the model
rnn_history = rnn_model.fit(X_train, y_train, epochs=20, batch_size=32, validation

# Evaluate the model
rnn_loss, rnn_accuracy = rnn_model.evaluate(X_test, y_test)
print("Test Accuracy:", rnn_accuracy)

```



RNN Model Architecture :
Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 31, 32)	851008
simple_rnn (SimpleRNN)	(None, 16)	784
dense_1 (Dense)	(None, 3)	51
activation (Activation)	(None, 3)	0

=====
Total params: 851843 (3.25 MB)
Trainable params: 851843 (3.25 MB)
Non-trainable params: 0 (0.00 Byte)

None

Epoch 1/20
687/687 [=====] - 35s 48ms/step - loss: 1.0884 - accu
Epoch 2/20
687/687 [=====] - 17s 25ms/step - loss: 1.0873 - accu
Epoch 3/20
687/687 [=====] - 24s 35ms/step - loss: 1.0692 - accu
Epoch 4/20
687/687 [=====] - 25s 36ms/step - loss: 1.0027 - accu
Epoch 5/20
687/687 [=====] - 33s 48ms/step - loss: 0.9185 - accu
Epoch 6/20
687/687 [=====] - 25s 37ms/step - loss: 0.8306 - accu
Epoch 7/20
687/687 [=====] - 23s 34ms/step - loss: 0.7526 - accu
Epoch 8/20
687/687 [=====] - 29s 42ms/step - loss: 0.6822 - accu

```

Epoch 9/20
687/687 [=====] - 30s 43ms/step - loss: 0.6137 - acc: 0.4500
Epoch 10/20
687/687 [=====] - 17s 25ms/step - loss: 0.5657 - acc: 0.4700
Epoch 11/20
687/687 [=====] - 16s 23ms/step - loss: 0.5245 - acc: 0.4900
Epoch 12/20
687/687 [=====] - 16s 24ms/step - loss: 0.4751 - acc: 0.5100
Epoch 13/20
687/687 [=====] - 16s 23ms/step - loss: 0.4464 - acc: 0.5300
Epoch 14/20
687/687 [=====] - 17s 25ms/step - loss: 0.4119 - acc: 0.5500
Epoch 15/20
687/687 [=====] - 17s 25ms/step - loss: 0.3788 - acc: 0.5700
Epoch 16/20
687/687 [=====] - 16s 23ms/step - loss: 0.3665 - acc: 0.5800
Epoch 17/20
687/687 [=====] - 16s 23ms/step - loss: 0.3392 - acc: 0.6000
Epoch 18/20
687/687 [=====] - 16s 24ms/step - loss: 0.3215 - acc: 0.6100
Epoch 19/20
687/687 [=====] - 16s 24ms/step - loss: 0.2991 - acc: 0.6300
Epoch 20/20
687/687 [=====] - 16s 24ms/step - loss: 0.2815 - acc: 0.6400

```

✓ Bi-Directional LSTM

```

from keras.layers import Bidirectional

# Build the Bidirectional LSTM model
bi_lstm_model = Sequential()
bi_lstm_model.add(Embedding(vocab_size, 32, input_length=max_len))
bi_lstm_model.add(Bidirectional(LSTM(64)))
bi_lstm_model.add(Dense(3, activation='softmax'))

bi_lstm_model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

print(" ")
print("Bi Directional LSTM Model Architecture : ")
print(bi_lstm_model.summary())
print(" ")

# Train the model
bi_lstm_history = bi_lstm_model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_val, y_val))

# Evaluate the model

```

```
bi_lstm_loss, bi_lstm_accuracy = bi_lstm_model.evaluate(X_test, y_test)
print("Test Accuracy:", bi_lstm_accuracy)
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 31, 32)	851008
bidirectional (Bidirectional)	(None, 128)	49664
dense_2 (Dense)	(None, 3)	387
Total params: 901059 (3.44 MB)		
Trainable params: 901059 (3.44 MB)		
Non-trainable params: 0 (0.00 Byte)		
None		
Epoch 1/20		
687/687 [=====] - 19s 23ms/step - loss: 0.8323 - accu		
Epoch 2/20		
687/687 [=====] - 7s 10ms/step - loss: 0.5654 - accu		
Epoch 3/20		
687/687 [=====] - 5s 8ms/step - loss: 0.4013 - accu		
Epoch 4/20		
687/687 [=====] - 6s 9ms/step - loss: 0.2981 - accu		
Epoch 5/20		
687/687 [=====] - 5s 8ms/step - loss: 0.2361 - accu		
Epoch 6/20		
687/687 [=====] - 6s 9ms/step - loss: 0.1918 - accu		
Epoch 7/20		
687/687 [=====] - 7s 10ms/step - loss: 0.1541 - accu		
Epoch 8/20		
687/687 [=====] - 5s 8ms/step - loss: 0.1277 - accu		
Epoch 9/20		
687/687 [=====] - 6s 9ms/step - loss: 0.1104 - accu		
Epoch 10/20		
687/687 [=====] - 5s 7ms/step - loss: 0.0888 - accu		
Epoch 11/20		
687/687 [=====] - 6s 9ms/step - loss: 0.0798 - accu		
Epoch 12/20		
687/687 [=====] - 5s 7ms/step - loss: 0.0675 - accu		
Epoch 13/20		
687/687 [=====] - 6s 9ms/step - loss: 0.0612 - accu		
Epoch 14/20		
687/687 [=====] - 5s 8ms/step - loss: 0.0531 - accu		
Epoch 15/20		


```

687/687 [=====] - 5s 7ms/step - loss: 0.0449 - accuracy: 0.6448
Epoch 16/20
687/687 [=====] - 6s 9ms/step - loss: 0.0431 - accuracy: 0.6448
Epoch 17/20
687/687 [=====] - 5s 7ms/step - loss: 0.0436 - accuracy: 0.6448
Epoch 18/20
687/687 [=====] - 6s 9ms/step - loss: 0.0330 - accuracy: 0.6448
Epoch 19/20
687/687 [=====] - 5s 7ms/step - loss: 0.0315 - accuracy: 0.6448
Epoch 20/20
687/687 [=====] - 6s 8ms/step - loss: 0.0270 - accuracy: 0.6448
111/111 [=====] - 0s 4ms/step - loss: 2.6877 - accuracy: 0.6448
Test Accuracy: 0.6448783278465271

```

✓ GRU Model

```

from keras.layers import GRU

# Build the GRU model
gru_model = Sequential()
gru_model.add(Embedding(vocab_size, 32, input_length=max_len))
gru_model.add(GRU(64))
gru_model.add(Dense(3, activation='softmax'))

gru_model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

print(" ")
print("Gated Recurrent Unit (GRU) Model Architecture : ")
print(gru_model.summary())
print(" ")

# Train the model
gru_history = gru_model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))

# Evaluate the model
gru_loss, gru_accuracy = gru_model.evaluate(X_test, y_test)
print("Test Accuracy:", gru_accuracy)

```



Gated Recurrent Unit (GRU) Model Architecture :
Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		

embedding_3 (Embedding)	(None, 31, 32)	851008
gru (GRU)	(None, 64)	18816
dense_3 (Dense)	(None, 3)	195

```

=====
Total params: 870019 (3.32 MB)
Trainable params: 870019 (3.32 MB)
Non-trainable params: 0 (0.00 Byte)

```

None

```

Epoch 1/20
687/687 [=====] - 16s 21ms/step - loss: 1.0882 - accu
Epoch 2/20
687/687 [=====] - 5s 7ms/step - loss: 1.0879 - accura
Epoch 3/20
687/687 [=====] - 4s 6ms/step - loss: 0.8354 - accura
Epoch 4/20
687/687 [=====] - 4s 5ms/step - loss: 0.5846 - accura
Epoch 5/20
687/687 [=====] - 4s 6ms/step - loss: 0.4039 - accura
Epoch 6/20
687/687 [=====] - 4s 6ms/step - loss: 0.2804 - accura
Epoch 7/20
687/687 [=====] - 4s 6ms/step - loss: 0.2043 - accura
Epoch 8/20
687/687 [=====] - 4s 5ms/step - loss: 0.1616 - accura
Epoch 9/20
687/687 [=====] - 5s 7ms/step - loss: 0.1301 - accura
Epoch 10/20
687/687 [=====] - 3s 5ms/step - loss: 0.1103 - accura
Epoch 11/20
687/687 [=====] - 3s 5ms/step - loss: 0.0942 - accura
Epoch 12/20
687/687 [=====] - 4s 6ms/step - loss: 0.0793 - accura
Epoch 13/20
687/687 [=====] - 4s 6ms/step - loss: 0.0715 - accura
Epoch 14/20
687/687 [=====] - 3s 5ms/step - loss: 0.0617 - accura
Epoch 15/20
687/687 [=====] - 4s 5ms/step - loss: 0.0543 - accura
Epoch 16/20
687/687 [=====] - 4s 6ms/step - loss: 0.0548 - accura
Epoch 17/20
687/687 [=====] - 4s 5ms/step - loss: 0.0487 - accura
Epoch 18/20
687/687 [=====] - 4s 5ms/step - loss: 0.0399 - accura
Epoch 19/20

```

```
687/687 [=====] - 4s 6ms/step - loss: 0.0380 - accuracy: 0.9500
Epoch 20/20
687/687 [=====] - 4s 6ms/step - loss: 0.0319 - accuracy: 0.9500
```

✓ CNN Model

```
from keras.layers import Conv1D, GlobalMaxPooling1D

# Build the CNN model
cnn_model = Sequential()
cnn_model.add(Embedding(vocab_size, 32, input_length=max_len))
cnn_model.add(Conv1D(128, 5, activation='relu'))
cnn_model.add(GlobalMaxPooling1D())
cnn_model.add(Dense(3, activation='softmax'))

cnn_model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

print(" ")
print(" Convolutional Neural Network (CNN) Model Architecture : ")
print(cnn_model.summary())
print(" ")

# Train the model
cnn_history = cnn_model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))

# Evaluate the model
cnn_loss, cnn_accuracy = cnn_model.evaluate(X_test, y_test)
print("Test Accuracy:", cnn_accuracy)
```



Convolutional Neural Network (CNN) Model Architecture :
Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 31, 32)	851008
conv1d (Conv1D)	(None, 27, 128)	20608
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense_4 (Dense)	(None, 3)	387

```
=====
Total params: 872003 (3.33 MB)
Trainable params: 872003 (3.33 MB)
Non-trainable params: 0 (0.00 Byte)
```

None

```
Epoch 1/20
687/687 [=====] - 16s 19ms/step - loss: 0.8165 - acc: 0.4125
Epoch 2/20
687/687 [=====] - 5s 7ms/step - loss: 0.5454 - acc: 0.5125
Epoch 3/20
687/687 [=====] - 3s 5ms/step - loss: 0.3268 - acc: 0.6125
Epoch 4/20
687/687 [=====] - 3s 4ms/step - loss: 0.1856 - acc: 0.7125
Epoch 5/20
687/687 [=====] - 4s 6ms/step - loss: 0.1115 - acc: 0.8125
Epoch 6/20
687/687 [=====] - 3s 4ms/step - loss: 0.0713 - acc: 0.9125
Epoch 7/20
687/687 [=====] - 4s 5ms/step - loss: 0.0482 - acc: 0.9625
Epoch 8/20
687/687 [=====] - 3s 4ms/step - loss: 0.0341 - acc: 0.9875
Epoch 9/20
687/687 [=====] - 4s 5ms/step - loss: 0.0248 - acc: 0.9975
Epoch 10/20
687/687 [=====] - 2s 4ms/step - loss: 0.0189 - acc: 1.0000
Epoch 11/20
687/687 [=====] - 3s 4ms/step - loss: 0.0170 - acc: 1.0000
Epoch 12/20
687/687 [=====] - 3s 4ms/step - loss: 0.0136 - acc: 1.0000
Epoch 13/20
687/687 [=====] - 3s 5ms/step - loss: 0.0130 - acc: 1.0000
Epoch 14/20
687/687 [=====] - 3s 4ms/step - loss: 0.0119 - acc: 1.0000
Epoch 15/20
687/687 [=====] - 3s 4ms/step - loss: 0.0118 - acc: 1.0000
Epoch 16/20
687/687 [=====] - 2s 4ms/step - loss: 0.0102 - acc: 1.0000
Epoch 17/20
687/687 [=====] - 2s 4ms/step - loss: 0.0095 - acc: 1.0000
Epoch 18/20
687/687 [=====] - 3s 5ms/step - loss: 0.0088 - acc: 1.0000
Epoch 19/20
```

✓ BERT Model

```

from transformers import BertTokenizer, TFBertForSequenceClassification
from keras.callbacks import EarlyStopping

# Load BERT tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

# Tokenize inputs
X_train_tokens = tokenizer(list(train_data['text']), padding=True, truncation=True)
X_test_tokens = tokenizer(list(test_data['text']), padding=True, truncation=True)

# Load pre-trained BERT model
bert_model = TFBertForSequenceClassification.from_pretrained('bert-base-uncased')

# Freeze BERT layers
for layer in bert_model.layers:
    layer.trainable = False

# Compile the model
bert_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])


print(" ")
print(" Bidirectional Encoder Representations from Transformers (BERT) Model Architecture")
print(bert_model.summary())
print(" ")


# Early stopping
early_stopping = EarlyStopping(monitor='val_accuracy', patience=2, restore_best_weights=True)

# Train the model
bert_history = bert_model.fit(X_train_tokens, y_train, epochs=10, batch_size=32,
                             validation_split=0.2, callbacks=[early_stopping])

# Evaluate the model
bert_loss, bert_accuracy = bert_model.evaluate(X_test_tokens, y_test)
print("Test Accuracy:", bert_accuracy)

```

 /usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:88: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab. You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models.
 warnings.warn(

tokenizer_config.json: 100%  48.0/48.0 [00:00<00:00, 2.43kB/s]

vocab.txt: 100%  232k/232k [00:00<00:00, 5.07MB/s]
 tokenizer.json: 100%  466k/466k [00:00<00:00, 3.68MB/s]
 config.json: 100%  570/570 [00:00<00:00, 42.1kB/s]
 model.safetensors: 100%  440M/440M [00:01<00:00, 227MB/s]
 All PyTorch model weights were used when initializing TFBertForSequenceClassif

Some weights or buffers of the TF 2.0 model TFBertForSequenceClassification we
 You should probably TRAIN this model on a down-stream task to be able to use i

Bidirectional Encoder Representations from Transformers (BERT) Model Architec
 Model: "tf_bert_for_sequence_classification"

Layer (type)	Output Shape	Param #
bert (TFBertMainLayer)	multiple	109482240
dropout_37 (Dropout)	multiple	0 (unused)
classifier (Dense)	multiple	1538

=====
 Total params: 109483778 (417.65 MB)
 Trainable params: 0 (0.00 Byte)
 Non-trainable params: 109483778 (417.65 MB)

None

Epoch 1/10

WARNING:tensorflow:AutoGraph could not transform <function infer_framework at
 Cause: for/else statement not yet supported

To silence this warning, decorate the function with @tf.autograph.experimental

WARNING: AutoGraph could not transform <function infer_framework at 0x7ee5e473

Cause: for/else statement not yet supported

To silence this warning, decorate the function with @tf.autograph.experimental

687/687 [=====] - 260s 316ms/step - loss: nan - accur

Epoch 2/10

687/687 [=====] - 202s 294ms/step - loss: nan - accur

Epoch 3/10

687/687 [=====] - 202s 295ms/step - loss: nan - accur

111/111 [=====] - 28s 154ms/step - loss: nan - accur

Test Accuracy: 0.4012450575828552

✓ Model Comparison

```
# Evaluate LSTM model
lstm_loss, lstm_accuracy = lstm_model.evaluate(X_test, y_test)
print("LSTM Test Accuracy:", lstm_accuracy)

# Evaluate Bidirectional LSTM model
bi_lstm_loss, bi_lstm_accuracy = bi_lstm_model.evaluate(X_test, y_test)
print("Bidirectional LSTM Test Accuracy:", bi_lstm_accuracy)

# Evaluate RNN model
rnn_loss, rnn_accuracy = rnn_model.evaluate(X_test, y_test)
print("RNN Test Accuracy:", rnn_accuracy)

# Evaluate GRU model
gru_loss, gru_accuracy = gru_model.evaluate(X_test, y_test)
print("GRU Test Accuracy:", gru_accuracy)

# Evaluate CNN model
cnn_loss, cnn_accuracy = cnn_model.evaluate(X_test, y_test)
print("CNN Test Accuracy:", cnn_accuracy)

# Evaluate BERT model
bert_loss, bert_accuracy = bert_model.evaluate(X_test_tokens, y_test)
print("BERT Test Accuracy:", bert_accuracy)
```

```
➡ 111/111 [=====] - 0s 3ms/step - loss: 2.1022 - accuracy: 0.6505376100540161
LSTM Test Accuracy: 0.6505376100540161
111/111 [=====] - 0s 4ms/step - loss: 2.6877 - accuracy: 0.6448783278465271
Bidirectional LSTM Test Accuracy: 0.6448783278465271
111/111 [=====] - 1s 5ms/step - loss: 1.7674 - accuracy: 0.4136955440044403
RNN Test Accuracy: 0.4136955440044403
111/111 [=====] - 0s 3ms/step - loss: 2.4522 - accuracy: 0.6511035561561584
GRU Test Accuracy: 0.6511035561561584
111/111 [=====] - 0s 2ms/step - loss: 2.5631 - accuracy: 0.6290322542190552
CNN Test Accuracy: 0.6290322542190552
111/111 [=====] - 17s 157ms/step - loss: nan - accuracy: 0.4012450575828552
BERT Test Accuracy: 0.4012450575828552
```

```
import pandas as pd

# Create a dictionary to store accuracy results
accuracy_results = {
    "Model": ["LSTM", "Bidirectional LSTM", "RNN", "GRU", "CNN", "BERT"],
    "Test Accuracy": [lstm_accuracy, bi_lstm_accuracy, rnn_accuracy, gru_accuracy,
}

# Convert the dictionary to a DataFrame
df_accuracy = pd.DataFrame(accuracy_results)

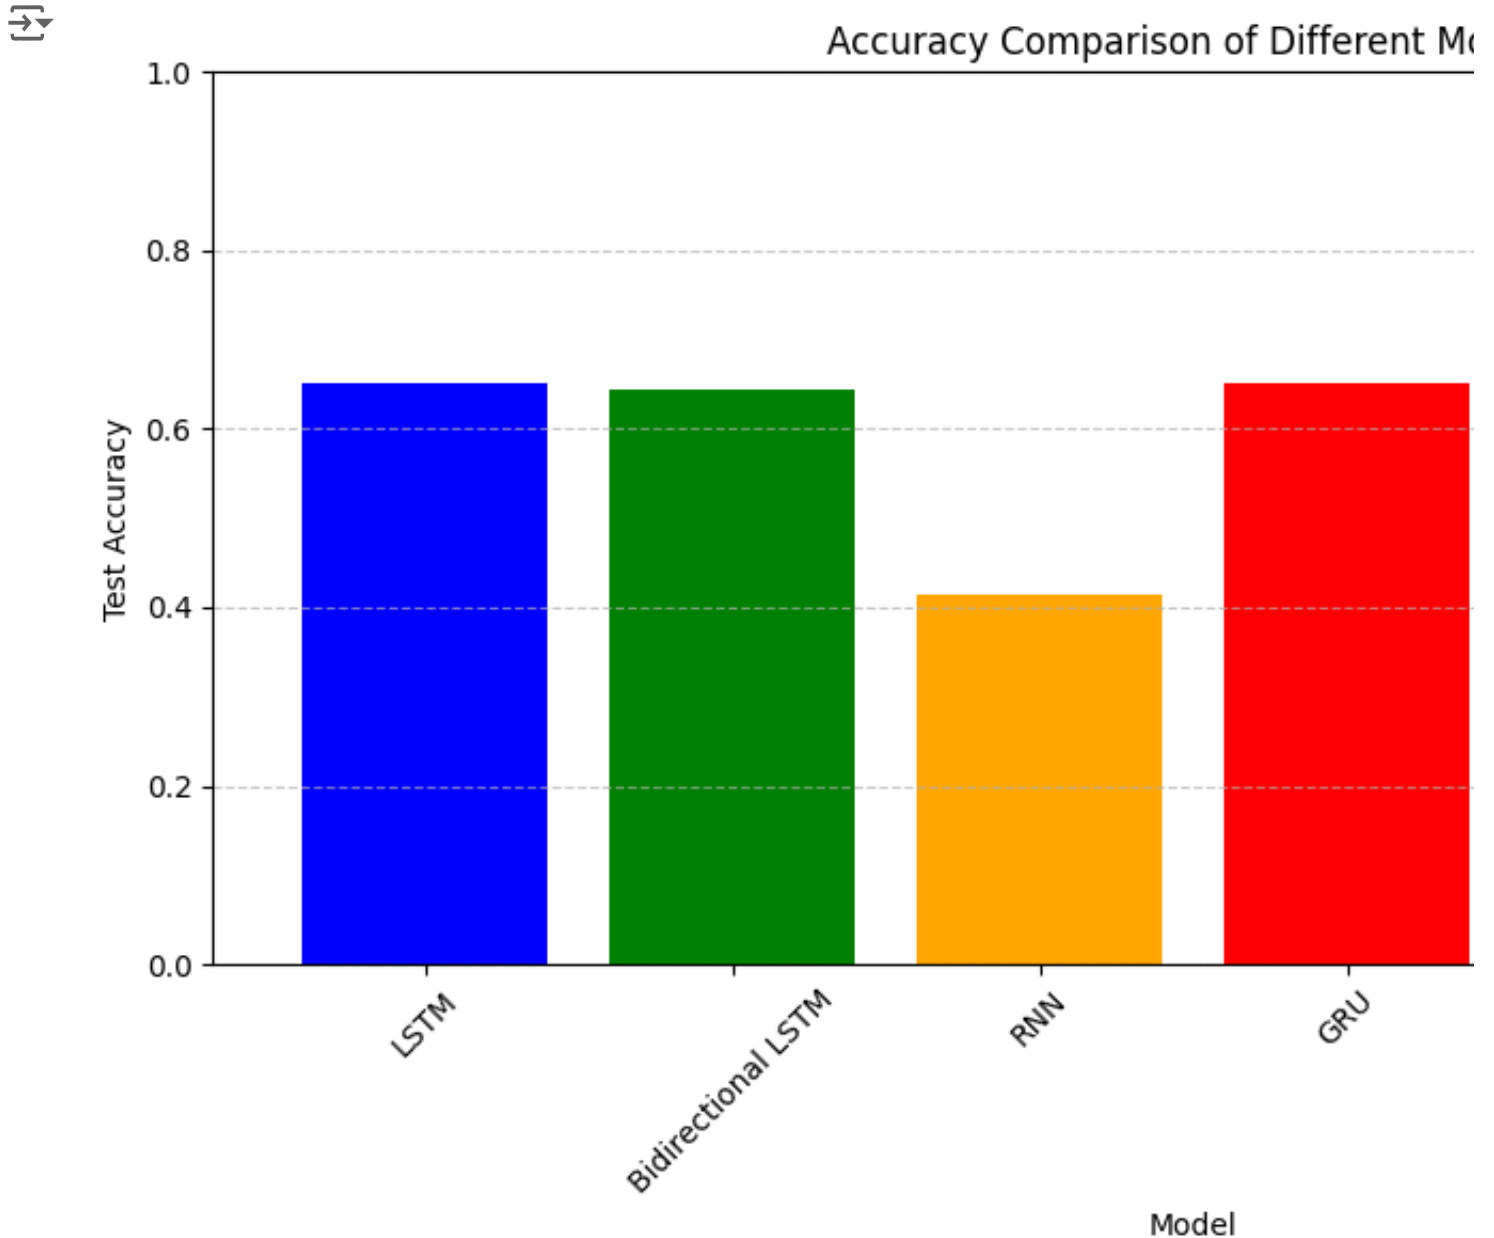
# Print the DataFrame
display(df_accuracy)
```



	Model	Test Accuracy
0	LSTM	0.650538
1	Bidirectional LSTM	0.644878
2	RNN	0.413696
3	GRU	0.651104
4	CNN	0.629032
5	BERT	0.401245

```
import matplotlib.pyplot as plt

# Plot accuracy results
plt.figure(figsize=(10, 6))
plt.bar(df_accuracy["Model"], df_accuracy["Test Accuracy"], color=['blue', 'green',
plt.xlabel('Model')
plt.ylabel('Test Accuracy')
plt.title('Accuracy Comparison of Different Models')
plt.ylim(0, 1) # Set y-axis limit from 0 to 1
plt.xticks(rotation=45) # Rotate x-axis labels for better visibility
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

```
import numpy as np
import seaborn as sns

# Make predictions for each model
lstm_preds = np.argmax(lstm_model.predict(X_test), axis=1)
bi_lstm_preds = np.argmax(bi_lstm_model.predict(X_test), axis=1)
rnn_preds = np.argmax(rnn_model.predict(X_test), axis=1)
gru_preds = np.argmax(gru_model.predict(X_test), axis=1)
cnn_preds = np.argmax(cnn_model.predict(X_test), axis=1)
bert_preds = np.argmax(bert_model.predict(X_test_tokens)[0], axis=1)

# Create a DataFrame to store the predictions
```

```
df_preds = pd.DataFrame({
    "LSTM": lstm_preds,
    "Bidirectional LSTM": bi_lstm_preds,
    "RNN": rnn_preds,
    "GRU": gru_preds,
    "CNN": cnn_preds,
    "BERT": bert_preds,
    "True Label": y_test
})


# Calculate the correlation matrix
correlation_matrix = df_preds.corr()

# Plot the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", square=True)
plt.title('Correlation Matrix of Model Predictions')
plt.show()
```

```

111/111 [=====] - 0s 3ms/step
111/111 [=====] - 0s 4ms/step
111/111 [=====] - 1s 6ms/step
111/111 [=====] - 0s 3ms/step
111/111 [=====] - 0s 3ms/step
111/111 [=====] - 17s 153ms/step

```



```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_sco
```

```

# Compute performance metrics for each model
metrics_dict = {
    "Model": ["LSTM", "Bidirectional LSTM", "RNN", "GRU", "CNN", "BERT"],
    "Accuracy": [accuracy_score(y_test, lstm_preds),
                  accuracy_score(y_test, bi_lstm_preds),
                  accuracy_score(y_test, rnn_preds),
                  accuracy_score(y_test, gru_preds),
                  accuracy_score(y_test, cnn_preds),
                  accuracy_score(y_test, bert_preds)],
    "Precision": [precision_score(y_test, lstm_preds, average='weighted'),
                  precision_score(y_test, bi_lstm_preds, average='weighted'),
                  precision_score(y_test, rnn_preds, average='weighted'),
                  precision_score(y_test, gru_preds, average='weighted'),
                  precision_score(y_test, cnn_preds, average='weighted'),
                  precision_score(y_test, bert_preds, average='weighted')],
    "Recall": [recall_score(y_test, lstm_preds, average='weighted'),
               recall_score(y_test, bi_lstm_preds, average='weighted'),
               recall_score(y_test, rnn_preds, average='weighted'),
               recall_score(y_test, gru_preds, average='weighted'),
               recall_score(y_test, cnn_preds, average='weighted'),
               recall_score(y_test, bert_preds, average='weighted')],
    "F1-score": [f1_score(y_test, lstm_preds, average='weighted'),
                  f1_score(y_test, bi_lstm_preds, average='weighted'),
                  f1_score(y_test, rnn_preds, average='weighted'),
                  f1_score(y_test, gru_preds, average='weighted'),
                  f1_score(y_test, cnn_preds, average='weighted'),
                  f1_score(y_test, bert_preds, average='weighted')]
}

# Create a DataFrame for the metrics
df_metrics = pd.DataFrame(metrics_dict)

# Plot the performance metrics
plt.figure(figsize=(12, 8))

# Accuracy
plt.subplot(2, 2, 1)
plt.bar(df_metrics["Model"], df_metrics["Accuracy"], color='blue')
plt.title('Accuracy of Different Models')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.xticks(rotation=45)

# Precision

```

```
plt.subplot(2, 2, 2)
plt.bar(df_metrics["Model"], df_metrics["Precision"], color='green')
plt.title('Precision of Different Models')
plt.xlabel('Model')
plt.ylabel('Precision')
plt.xticks(rotation=45)

# Recall
plt.subplot(2, 2, 3)
plt.bar(df_metrics["Model"], df_metrics["Recall"], color='orange')
plt.title('Recall of Different Models')
plt.xlabel('Model')
plt.ylabel('Recall')
plt.xticks(rotation=45)

# F1-score
plt.subplot(2, 2, 4)
plt.bar(df_metrics["Model"], df_metrics["F1-score"], color='red')
plt.title('F1-score of Different Models')
plt.xlabel('Model')
plt.ylabel('F1-score')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

```
→ /usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:134
   _warn_prf(average, modifier, msg_start, len(result))
```

